
GPU FFT Filtering of Seismic Data

Capstone Project

Challenges

- Writing code from scratch to read and write SEG-Y format seismic data
- Learning how to use cuFFT to do a 1-D R2C and then C2R transform, although it appeared to work, it was not until adding in cuFFT error checking to find that one of the cuFFT calls was silently failing
- Learning about Nyquist and Frequencies so I could properly write the kernel for filter the data samples based on a range of minimum and maximum frequency to keep
- All class example code generally centered around filtering pictures, so doing a project like this meant you had to figure things out for yourself



1. Intro

Seismic Data is the recording of sound/vibration by geophones when a source like dynamite or vibrators is set off. Filtering the data removing high and low frequencies can help to remove unwanted noise from the data.

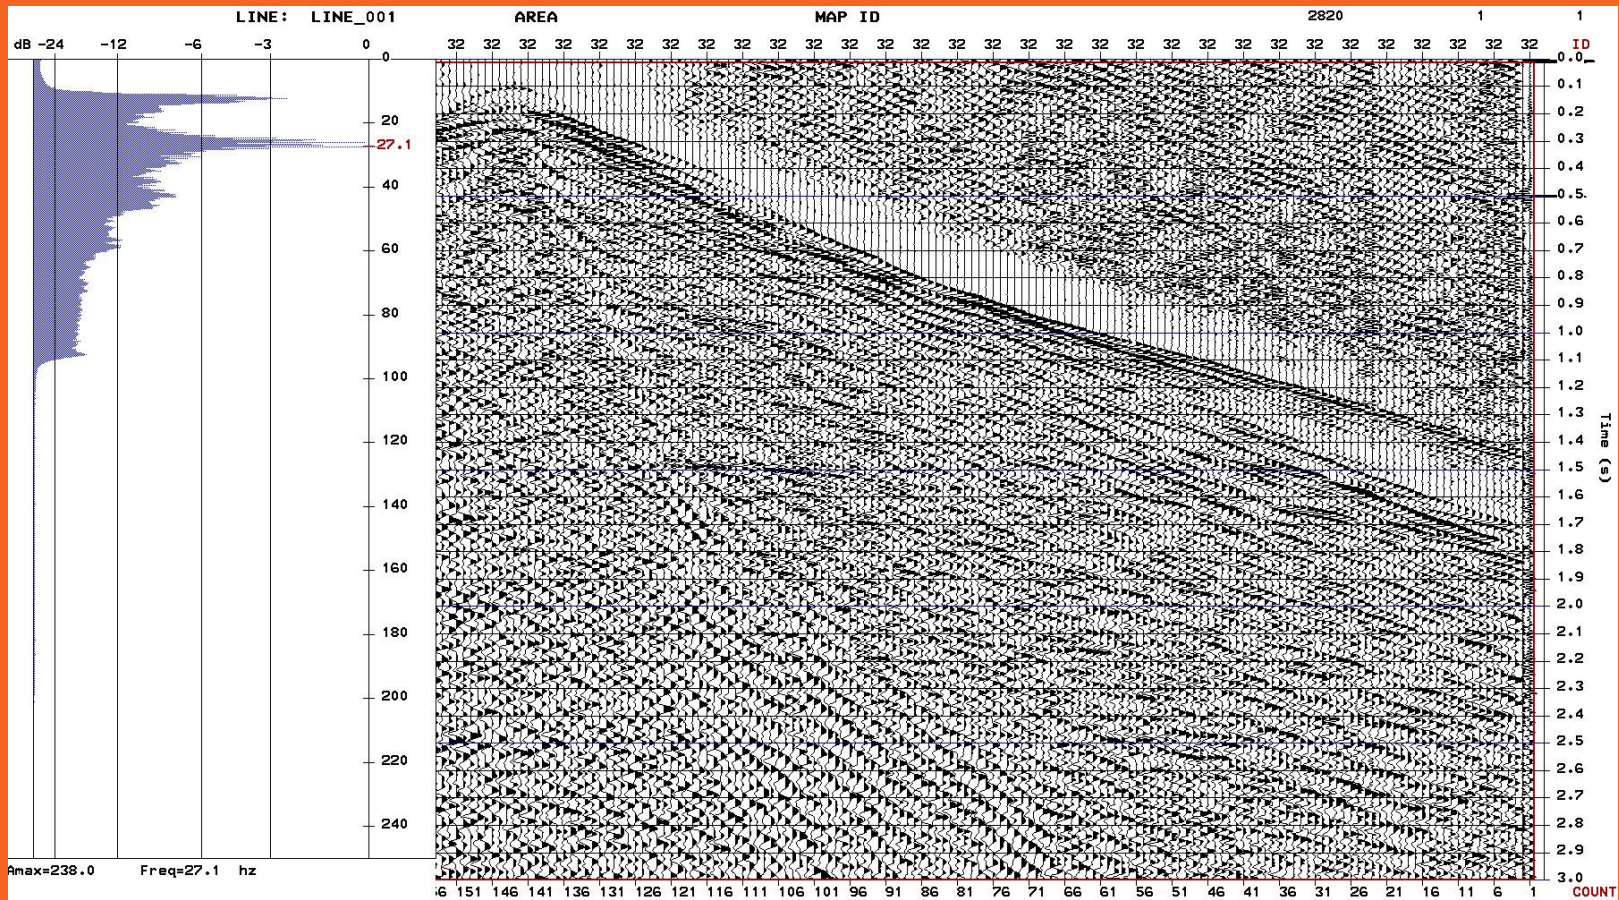
Fast Fourier Transform

Converts from time domain to frequency domain, allowing you to work on specific frequencies in hertz


Nyquist

Nyquist frequency (*cycles per second*) is the frequency whose cycle-length (or period) is twice the interval between samples, thus *0.5 cycle/sample*.

This information along with sample rate and number of samples is required to figure out what frequencies are contained within the converted FFT buffers the GPU memory.



Input Data before filtering, left hand graph shows frequency content through most of the range from 0 to 250 HZ



Output data after running with 30 to 60 min max parameters showing data concentrated in the range from 30 to 60 HZ

30 to 60 Run log:

CUDA Device name: NVIDIA GeForce GTX 1650 Ti with Max-Q
Design

GPU Count: (1)

Core Count: (16)

ARCHITECTURE -> little endian (0)

Input File: ../data/Line_1_IEEE_Float_LE.sgy

Output File: output_30_60.sgy

Minimum Frequency: 30

Maximum Frequency: 60

Simultaneous Traces: 10000

Open for read: ../data/Line_1_IEEE_Float_LE.sgy

EBCDIC Header Read 3200 Bytes

Binary Header Read 400 Bytes

Ebcdic Header: EBCDIC

Revision < 2.0 Detected

Binary Header: IBM ORDER

Num Traces: -31073 (FFFF869F) bytes 3213-3214

Num Aux Tr: 0 (0000) bytes 3215-3216

Sample Interval: 2000 (07D0) bytes 3217-3218

Num Samples: 1501 (05DD) bytes 3221-3222

Revision: (0) 0 (0000) bytes 3297-3300

Format: 11 (000B) bytes 3225-3226

Calculated Number of Traces: 2820

Open for write: output_30_60.sgy

First/Last 8 data samples for QC:

Input 0: 4.49441051

Input 1: 6.59748077

Input 2: 1.05115414

Input 3: -5.72130013

Input 4: -5.60891533

Input 5: 1.45248413

Input 6: 7.15355396

Input 7: 4.57948112

Input 15009992: 0.00000000

Input 15009993: 0.00000000

Input 15009994: 0.00000000

Input 15009995: 0.00000000

Input 15009996: 0.00000000

Input 15009997: 0.00000000

Input 15009998: 0.00000000

Input 15009999: 0.00000000

gufftflt_32f_i (60040000) 17611680/17611680

Output 0: 20292.77539062

Output 1: 7613.83007812

Output 2: -3015.89257812

Output 3: -4689.10351562

Output 4: 3657.17187500

Output 5: 16573.44921875

Output 6: 25208.56835938

Output 7: 22326.87304688

Output 15009992: 0.00000000

Output 15009993: 0.00000000

Output 15009994: 0.00000000

Output 15009995: 0.00000000

Output 15009996: 0.00000000

Output 15009997: 0.00000000

Output 15009998: 0.00000000

Output 15009999: 0.00000000

Total Traces Processed: 2820

Total Samples Processed (float numbers): 4232820

Elapsed Time: 46590706 microseconds

50 to 100 Run log:

CUDA Device name: NVIDIA GeForce GTX 1650 Ti with Max-Q Design

GPU Count: (1)

Core Count: (16)

ARCHITECTURE -> little endian (0)

Input File: ../data/Line_1_IEEE_Float_LE.sgy

Output File: output_50_100.sgy

Minimum Frequency: 50

Maximum Frequency: 100

Simultaneous Traces: 10000

Open for read: ../data/Line_1_IEEE_Float_LE.sgy

EBCDIC Header Read 3200 Bytes

Binary Header Read 400 Bytes

Ebcdic Header: EBCDIC

Revision < 2.0 Detected

Binary Header: IBM ORDER

Num Traces: -31073 (FFFF869F) bytes 3213-3214

Num Aux Tr: 0 (0000) bytes 3215-3216

Sample Interval: 2000 (07D0) bytes 3217-3218

Num Samples: 1501 (05DD) bytes 3221-3222

Revision: (0) 0 (0000) bytes 3297-3300

Format: 11 (000B) bytes 3225-3226

Calculated Number of Traces: 2820

Open for write: output_50_100.sgy

First/Last 8 data samples for QC:

Input 0: 4.49441051

Input 1: 6.59748077

Input 2: 1.05115414

Input 3: -5.72130013

Input 4: -5.60891533

Input 5: 1.45248413

Input 6: 7.15355396

Input 7: 4.57948112

Input 15009992: 0.00000000

Input 15009993: 0.00000000

Input 15009994: 0.00000000

Input 15009995: 0.00000000

Input 15009996: 0.00000000

Input 15009997: 0.00000000

Input 15009998: 0.00000000

Input 15009999: 0.00000000

gufftflt_32f_i (60040000) 17611680/17611680

Output 0: 19214.82031250

Output 1: 27434.47851562

Output 2: 20398.21093750

Output 3: 5081.96093750

Output 4: -6893.46484375

Output 5: -10853.37500000

Output 6: -11840.56250000

Output 7: -15484.06640625

Output 15009992: 0.00000000

Output 15009993: 0.00000000

Output 15009994: 0.00000000

Output 15009995: 0.00000000

Output 15009996: 0.00000000

Output 15009997: 0.00000000

Output 15009998: 0.00000000

Output 15009999: 0.00000000

Total Traces Processed: 2820

Total Samples Processed (float numbers): 4232820

Elapsed Time: 47010962 microseconds



2. Results

Filtered Data

Output data was correctly filtered by user supplied minimum and maximum frequency

Repeatability

The data was run and compared multiple times to ensure there was luck involved and the transform and filter worked correctly

Robustness

The program was run on data sets over 500MB containing millions of data samples and functioned correctly. Unfortunately these data sets could not be uploaded to GitHub as they have size limit restrictions



3. Learning

Troubles I ran into and things I learned :

FFT Buffers

Learned why you need to divide by 2 and then add 1 for cuFFT buffers, Realized how your time samples are now half the amount but contain a real/imaginary part representing amplitude and phase.

NVtop

The importance of trying large data sets and using NVtop to troubleshoot. Upon a large run the program reported memory inconsistencies, NVtop showed a growing memory usage while running. I tracked it down to the cuFFT plan not being freed. A small run did not crash, so use larger data sets and profile your run with NVtop to ensure your run is sustainable. Free you buffers and plans!



4. More Learning

Even more things I learned :

Cuda Functions

Unless there is an actual fault, cuda functions will happily return quietly while doing nothing they were supposed to. You must write all error checking of functions from the very start, or you will chase problems needlessly for hours.

Parameters

Test many different parameters and data sets, they will let you find bugs you did not see and let you find performance improvements. Also test your code with parameters that do not change the data, this will ensure that you are able to reproduce your input correctly, before you start adding complex filtering.